

Assignment 2: Clustering and Kernels

Members: Cui(Cathy) Yu, Shravani Mahadeshwar, Priya Botsa, Marc Hayes, Shaonan Feng and Di Yang

- *Course: Machine Learning 2*
- *Prof: Jan Nagler*
- *Due: 25 April 2026*

Methodology

Part 1: Clustering

Dataset

- Seeds dataset (UCI, ID 236) - 210 wheat kernel samples across 3 varieties: **Kama, Rosa, and Canadian**
- Chosen because the three classes form partially overlapping, roughly elliptical clusters, well-suited to compare a centroid-based method (*K-Means*) against a probabilistic one (*GMM*) and to demonstrate where a density-based method (*DBSCAN*) breaks down on uniformly distributed data.

Clustering Algorithms

- **K-Means:** partitions data by minimising within-cluster inertia; assumes spherical, equal-size clusters; hard assignment
- **GMM (full covariance):** models each cluster as a multivariate Gaussian; supports elliptical boundaries and soft probabilistic membership; evaluated on BIC and AIC for model fit
- **DBSCAN:** density-based method that groups points reachable within radius `eps` and flags low-density regions as noise; makes no assumptions about cluster shape or count; included as a contrast case to demonstrate where density-based approaches break down on uniformly distributed data

Part 2: Kernel SVM

Dataset

- Heart Disease dataset (UCI, ID 45) - 303 patient records, 13 clinical features (age, cholesterol, resting BP, chest pain type, etc.)
- Binary target: **0 = no heart disease, 1 = disease** (derived from original multi-class "num" column via thresholding at > 0)
- Chosen because clinical features are expected to have nonlinear interactions, making it a meaningful testbed for kernel comparison

Metrics

- Two metrics were chosen to reflect opposing clinical priorities, using the general F-beta formula:

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

- The **F2 score** ($\beta=2$) weights recall twice as much as precision, prioritising the correct identification of patients with heart disease. This is critical in medical diagnosis where false negatives (missed diagnoses) carry the highest risk.
- The **F0.5 score** ($\beta=0.5$) emphasises precision over recall, reducing false positives. This is relevant to avoid unnecessary stress, follow-up procedures, and medical costs for healthy patients.
- By optimizing both metrics, we evaluate model performance under two different practical priorities, reflecting real-world trade-offs in clinical decision-making.

Kernels and Hyperparameter Search

- Three kernels evaluated: **Linear (baseline), RBF, Polynomial**
- Linear kernel included as a baseline to assess whether nonlinear modeling adds meaningful improvement over a simpler decision boundary
- GridSearchCV with 5-fold cross-validation used for all searches
- Linear: $C \in \{0.01, 0.1, 1, 10, 100\}$
- RBF: $C \in \{0.01, 0.1, 1, 10, 100\}$, $\gamma \in \{0.001, 0.01, 0.1, 1\}$
- Polynomial: $C \in \{0.01, 0.1, 1, 10, 100\}$, $\text{degree} \in \{2, 3, 4, 5\}$
- Each kernel optimised separately for **F2** and **F0.5**, yielding 6 trained models total

1. Cluster the Clustering Monster

Importing necessary libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import cluster
import matplotlib.cm as cm
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
from scipy.optimize import linear_sum_assignment
from sklearn.mixture import GaussianMixture
from matplotlib.patches import Ellipse
import warnings
from sklearn.cluster import DBSCAN
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import NearestNeighbors
from sklearn.metrics import fbeta_score, make_scorer, classification_report
```

```
warnings.filterwarnings("ignore")
np.random.seed(42)
```

Loading the dataset

```
In [2]: col_names = ["area", "perimeter", "compactness", "length_kernel", "width_kernel",
                    "asymmetry_coeff", "length_groove", "target",]
df = pd.read_csv("data/seeds_dataset.txt", sep="\s+", header=None, names=col_names)
y_true = df["target"]
X_raw = df.drop(columns=["target"])
print("Shape:", X_raw.shape)
print("Classes:", y_true.unique())
X_raw.head()
```

Shape: (210, 7)

Classes: [1 2 3]

```
Out[2]:
```

	area	perimeter	compactness	length_kernel	width_kernel	asymmetry_coeff	length_groove
0	15.26	14.84	0.8710	5.763	3.312	2.221	5.220
1	14.88	14.57	0.8811	5.554	3.333	1.018	4.956
2	14.29	14.09	0.9050	5.291	3.337	2.699	4.825
3	13.84	13.94	0.8955	5.324	3.379	2.259	4.805
4	16.14	14.99	0.9034	5.658	3.562	1.355	5.175

```
In [3]: print("Missing values:", X_raw.isnull().sum().sum())
```

Missing values: 0

Scaling the features

```
In [4]: scaler = StandardScaler()
X = scaler.fit_transform(X_raw)
print("Scaled shape:", X.shape)
```

Scaled shape: (210, 7)

Plotting the first two principal components to visualize the data distribution


```
In [5]: pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)
var_explained = pca.explained_variance_ratio_.sum()
print(f"Variance explained by PC1+PC2: {var_explained:.2%}")
fig, ax = plt.subplots(figsize=(8, 6))
```

```

ax.scatter(X_pca[:, 0], X_pca[:, 1], s=60, alpha=0.7, color="#4C72B0", edgecolors="white", linewidths=0.5)
ax.set_title("PCA", fontsize=14, fontweight="bold", pad=15)
ax.set_xlabel(f"PC1 ({pca.explained_variance_ratio_[0]:.1%} variance)", fontsize=12)
ax.set_ylabel(f"PC2 ({pca.explained_variance_ratio_[1]:.1%} variance)", fontsize=12)
ax.grid(True, linestyle="--", alpha=0.4)
fig.tight_layout()
plt.show()

```

Variance explained by PC1+PC2: 88.98%



- PC1 dominates with 71.9% variance, meaning most meaningful differences between wheat kernels come down to a single dimension likely overall **kernel size**.

- The 2D PCA projection already captures ~89% of the total data variance, so very little information is lost by reducing from 7 features to 2 components (for visualising the data).

Elbow and Silhouette: Plotting inertia and silhouette scores for K=2 to 10 for hyperparameter tuning

```
In [6]: inertias = []
K_range = range(2, 11)
for k in K_range:
    km = cluster.KMeans(n_clusters=k, random_state=42, n_init=10)
    km.fit(X)
    inertias.append(km.inertia_)


fig, ax = plt.subplots(1, 1, figsize=(7, 5))
fig.suptitle("K-Means Hyperparameter Tuning", fontsize=14, fontweight="bold", y=1.02)

ax.plot(K_range, inertias, "o-", color="#2196F3", linewidth=2.5, markersize=8, markerfacecolor="white", markeredgewidth=2)
ax.axvline(x=3, color="red", linestyle="--", alpha=0.6, label="K=3 (elbow)")
ax.set_title("Elbow Method", fontsize=13, fontweight="bold")
ax.set_xlabel("Number of Clusters K", fontsize=11)
ax.set_ylabel("Inertia", fontsize=11)
ax.legend(fontsize=10)
ax.grid(True, linestyle="--", alpha=0.4)

best_k = 3
print(f"Chosen K = {best_k}")
```

Chosen K = 3

K-Means Hyperparameter Tuning



Parameter Interpretation:

- **K=3** is the best choice as the elbow plot shows a **sharp drop in inertia from K=2 to K=3**, after which the curve flattens significantly, indicating that adding more clusters yields diminishing returns.

```
In [7]: # Creating a mapping function to ensure consistent cluster labels across K-Means and GMM
def align_labels(ref, new, n):
    """Aligning new cluster labels to reference labels via Hungarian algorithm."""
    cost = np.zeros((n, n))
    for i in range(n):
        for j in range(n):
            cost[i, j] = -np.sum((ref == i) & (new == j))
    _, col_ind = linear_sum_assignment(cost)
    return col_ind, col_ind[new]
```

```
def draw_ellipse(mean, cov, ax, color, pca, alpha=0.15, n_std=2.0):
    """Drawing a confidence ellipse for a GMM component in PCA space."""
    cov_pca = pca.components_ @ cov @ pca.components_.T
    vals, vecs = np.linalg.eigh(cov_pca)
    angle = np.degrees(np.arctan2(*vecs[:, -1][::-1]))
    width, height = 2 * n_std * np.sqrt(vals[::-1])
    ax.add_patch(
        Ellipse(xy=mean, width=width, height=height, angle=angle, color=color, alpha=alpha))
    ax.add_patch(
        Ellipse(xy=mean, width=width, height=height, angle=angle, fill=False, edgecolor=color, linewidth=2, linestyle="--"))
palette = ["#1F9EF8FF", "#F063B5", "#FDDF48"]
```

KMeans

```
In [8]: kmeans = cluster.KMeans(n_clusters=best_k, random_state=42, n_init=10)
kmeans.fit(X)
labels_km = kmeans.labels_
# Projecting K-Means centers to PCA space for visualization
centers_pca = pca.transform(kmeans.cluster_centers_)
```

Gaussian Mixture Models

```
In [9]: gmm = GaussianMixture(n_components=3, covariance_type="full", random_state=42)
gmm.fit(X)
labels_gmm = gmm.predict(X)
probs = gmm.predict_proba(X)
print(f"GMM converged: {gmm.converged_}")
print(f"GMM BIC: {gmm.bic(X):.2f}")
print(f"GMM AIC: {gmm.aic(X):.2f}")
col_ind, labels_gmm_aligned = align_labels(labels_km, labels_gmm, best_k)
gmm_means_aligned = gmm.means_[col_ind]
gmm_covs_aligned = gmm.covariances_[col_ind]
X_df = X_raw.copy()
X_df["cluster"] = labels_km
X_df.groupby("cluster").mean().round(3)
# Projecting GMM means to PCA space for visualization
means_pca = pca.transform(gmm_means_aligned)
```

GMM converged: True


GMM BIC: -69.70

GMM AIC: -427.84

- GMM converged successfully; the negative BIC (-69.70) and AIC (-427.84) confirm the model is a good fit, with lower (more negative) values indicating better fit relative to model complexity.

DBSCAN

```
In [10]: # Finding k-distance for k=min_samples=5 for DBSCAN eps tuning
nbrs = NearestNeighbors(n_neighbors=5).fit(X)
distances, _ = nbrs.kneighbors(X)
k_distances = np.sort(distances[:, -1])[:, -1]
plt.figure(figsize=(8, 4))
plt.plot(k_distances, color="#2196F3", linewidth=2)
plt.axhline(y=1.0, color="red", linestyle="--", alpha=0.6, label="eps=1.0")
plt.title("K-Distance Plot (k=5)", fontsize=13, fontweight="bold")
plt.xlabel("Points sorted by distance", fontsize=11)
plt.ylabel("5th Nearest Neighbor Distance", fontsize=11)
plt.legend()
plt.grid(True, linestyle="--", alpha=0.4)
plt.tight_layout()
plt.show()
```



- The k-distance plot shows **no sharp elbow**, indicating the Seeds data has uniform density with no natural gaps, which is a fundamental mismatch with DBSCAN's assumptions.
- DBSCAN requires density contrasts to separate clusters, which this dataset does not provide.
- Therefore DBSCAN is fundamentally mismatched to this data geometry, and its failure is expected regardless of eps.

```
In [11]: db = DBSCAN(eps=1.0, min_samples=5)
db_labels = db.fit_predict(X)
n_clusters_db = len(set(db_labels) - {-1})
n_noise = np.sum(db_labels == -1)
print(f"DBSCAN clusters found: {n_clusters_db}")
print(f"Noise points: {n_noise}")
```

DBSCAN clusters found: 1

Noise points: 13

Comparing clustering algos

```
In [12]: unique_db_labels = sorted(set(db_labels))
n_db_clusters = len([l for l in unique_db_labels if l != -1])
db_colors = cm.tab10(np.linspace(0, 0.9, max(n_db_clusters, 1)))
db_palette = {lbl: db_colors[i] for i, lbl in enumerate(l for l in unique_db_labels if l != -1)}
# Visualizing K-Means, GMM, and DBSCAN results side by side in PCA space
fig, axes = plt.subplots(1, 3, figsize=(22, 7))
fig.suptitle("K-Means vs GMM vs DBSCAN", fontsize=15, fontweight="bold", y=1.02)
for i in range(best_k):
    mask_km = labels_km == i
    axes[0].scatter(X_pca[mask_km, 0], X_pca[mask_km, 1], s=60, alpha=0.85,
                    color=palette[i], edgecolors="white", linewidths=0.5,
                    label=f"Cluster {i}")
    axes[0].scatter(centers_pca[i, 0], centers_pca[i, 1], s=220, color=palette[i],
                    marker="X", edgecolors="black", linewidths=1.5, zorder=5)
    mask_gmm = labels_gmm_aligned == i
    axes[1].scatter(X_pca[mask_gmm, 0], X_pca[mask_gmm, 1], s=60, alpha=0.85,
                    color=palette[i], edgecolors="white", linewidths=0.5,
                    label=f"Component {i} ")
    draw_ellipse(means_pca[i], gmm_covs_aligned[i], axes[1], palette[i], pca, alpha=0.12)
    axes[1].scatter(means_pca[i, 0], means_pca[i, 1], s=220, color=palette[i],
                    marker="X", edgecolors="black", linewidths=1.2, zorder=5)
# DBSCAN
for lbl in unique_db_labels:
    mask_db = db_labels == lbl
    if lbl == -1:
        axes[2].scatter(X_pca[mask_db, 0], X_pca[mask_db, 1], s=40, alpha=0.4,
                        color="#999999", marker="x", label="Noise")
    else:
```


```

axes[2].scatter(X_pca[mask_db, 0], X_pca[mask_db, 1], s=60, alpha=0.85,
               color=db_palette[lbl], edgecolors="white", linewidths=0.5,
               label=f"Cluster {lbl}")
axes[2].text(0.02, 0.97, f"Clusters: {n_db_clusters} | Noise: {n_noise} pts",
            transform=axes[2].transAxes, fontsize=9, va="top", color="#444444")

for ax, title in zip(axes, [f"K-Means (K={best_k}): Hard Clustering",
                             "Gaussian Mixture Model: Soft Probabilistic Clusters",
                             f"DBSCAN (eps={db.eps}, min_samples={db.min_samples}): Hard Clusters + Noise"]):
    ax.set_title(title, fontsize=13, fontweight="bold")
    ax.set_xlabel(f"PC1 ({{pca.explained_variance_ratio_[0]:.1%}})", fontsize=11)
    ax.set_ylabel(f"PC2 ({{pca.explained_variance_ratio_[1]:.1%}})", fontsize=11)
    ax.legend(fontsize=10)
    ax.grid(True, linestyle="--", alpha=0.4)
fig.tight_layout()
plt.show()

```

K-Means vs GMM vs DBSCAN



Clustering Performance Evaluation using Silhouette Score

```

In [13]: print("KMeans silhouette:", round(silhouette_score(X, labels_km), 4))
          print("GMM silhouette   :", round(silhouette_score(X, labels_gmm), 4))
          # Calculating silhouette for DBSCAN excluding noise points (-1 labels)
          mask_valid = db_labels != -1
          if mask_valid.sum() > 1 and len(set(db_labels[mask_valid])) > 1:
              db_sil = round(silhouette_score(X[mask_valid], db_labels[mask_valid]), 4)
              print(f"DBSCAN silhouette (excl. noise): {db_sil}")
          else:

```

```
print(f"DBSCAN silhouette: undefined (only {len(set(db_labels[mask_valid]))} cluster found)")  
# print("DBSCAN silhouette (with noise):", round(silhouette_score(X, db_labels), 4))
```

KMeans silhouette: 0.4007

GMM silhouette : 0.3625

DBSCAN silhouette: undefined (only 1 cluster found)

Conclusion

- **K-Means** scores highest (0.40) as clusters are compact and well-separated, confirming it is the best fit for this dataset
- **Gaussian Mixture Model** scores slightly lower (0.36) with soft boundaries introducing minor overlap between components, but still recovers the 3 varieties reliably
- **DBSCAN** produces an undefined silhouette score as it collapses all points into a single cluster, failing to detect any structure due to the absence of density gaps in the data.

Overall, **K-Means** is the recommended algorithm for the Seeds dataset as it achieves the **highest silhouette score, recovers all 3 wheat varieties correctly**, and its assumption of spherical evenly-sized clusters aligns naturally with the true underlying class structure

2. Kernels

Loading the dataset

```
In [14]: from ucimlrepo import fetch_ucirepo  
heart_disease = fetch_ucirepo(id=45)
```

```
In [15]: # metadata of the dataset  
# print(heart_disease.metadata)  
X = heart_disease.data.features  
y = heart_disease.data.targets  
# display variable information to understand the features and target  
heart_disease.variables
```

Out[15]:

	name	role	type	demographic	description	units	missing_values	
0	age	Feature	Integer	Age		NaN	years	no
1	sex	Feature	Categorical	Sex		NaN	NaN	no
2	cp	Feature	Categorical	NaN		NaN	NaN	no
3	trestbps	Feature	Integer	NaN	resting blood pressure (on admission to the ho...	mm Hg		no
4	chol	Feature	Integer	NaN	serum cholestoral	mg/dl		no
5	fbs	Feature	Categorical	NaN	fasting blood sugar > 120 mg/dl	NaN		no
6	restecg	Feature	Categorical	NaN		NaN	NaN	no
7	thalach	Feature	Integer	NaN	maximum heart rate achieved	NaN		no
8	exang	Feature	Categorical	NaN	exercise induced angina	NaN		no
9	oldpeak	Feature	Integer	NaN	ST depression induced by exercise relative to ...	NaN		no
10	slope	Feature	Categorical	NaN		NaN	NaN	no
11	ca	Feature	Integer	NaN	number of major vessels (0-3) colored by flour...	NaN		yes
12	thal	Feature	Categorical	NaN		NaN	NaN	yes
13	num	Target	Integer	NaN	diagnosis of heart disease	NaN		no

Pre-processing the data

```
In [16]: # Drop rows with missing values
X_hd = heart_disease.data.features.copy()
y_hd = heart_disease.data.targets.copy()
X_hd = X_hd.dropna()
y_hd = y_hd.loc[X_hd.index]
# Binary classification: 0 = no disease, 1 = disease (num > 0)
y_binary = (y_hd["num"] > 0).astype(int)
# Scale features
scaler_hd = StandardScaler()
X_scaled = scaler_hd.fit_transform(X_hd)
# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_binary, test_size=0.2, random_state=42, stratify=y_binary)
print("Train size:", X_train.shape, " Test size:", X_test.shape)
print("Class distribution (train):", y_train.value_counts().to_dict())
```

```
Train size: (237, 13) Test size: (60, 13)
Class distribution (train): {0: 128, 1: 109}
```

Defining f2_scorer and f05_scorer

```
In [17]: f2_scorer = make_scorer(fbeta_score, beta=2)
f05_scorer = make_scorer(fbeta_score, beta=0.5)
print("F2 scorer:", f2_scorer)
print("F0.5 scorer:", f05_scorer)
```

F2 scorer: make_scorer(fbeta_score, response_method='predict', beta=2)

F0.5 scorer: make_scorer(fbeta_score, response_method='predict', beta=0.5)

- We expect the relationship between clinical features and heart disease **not to be linearly separable**, so *RBF and polynomial kernels* are used to capture complex decision boundaries.
- We include a **linear kernel** as a baseline to assess whether nonlinear modeling adds meaningful improvement.

Baseline SVC (linear kernel)

- **C** controls the trade-off between a wider margin and fewer misclassifications.

High C gives a stricter boundary while low C gives a softer margin.

```
In [18]: # Hyperparameter grid for linear SVC
param_grid_linear = {"C": [0.01, 0.1, 1, 10, 100]}
# optimize for F2 score
gs_linear_f2 = GridSearchCV(
    SVC(kernel="linear", random_state=42),
    param_grid_linear,
    cv=5,
    scoring=f2_scorer,
    return_train_score=False,
)
gs_linear_f2.fit(X_train, y_train)
best_linear_f2 = gs_linear_f2.best_estimator_
y_pred_linear_f2 = best_linear_f2.predict(X_test)
print("=== Linear SVC optimized for F2 ===")
print("Best params (linear):", gs_linear_f2.best_params_)
print("Test F2 :", round(fbeta_score(y_test, y_pred_linear_f2, beta=2), 4))
print("Test F0.5:", round(fbeta_score(y_test, y_pred_linear_f2, beta=0.5), 4))
print(
    "\n",
    classification_report(
        y_test, y_pred_linear_f2, target_names=["No Disease", "Disease"]),)
```

```

=== Linear SVC optimized for F2 ===
Best params (linear): {'C': 10}
Test F2 : 0.8029
Test F0.5: 0.8594

```

	precision	recall	f1-score	support
No Disease	0.83	0.91	0.87	32
Disease	0.88	0.79	0.83	28
accuracy			0.85	60
macro avg	0.85	0.85	0.85	60
weighted avg	0.85	0.85	0.85	60

```

In [19]: # optimize for F0.5 score
gs_linear_f05 = GridSearchCV(
    SVC(kernel="linear", random_state=42),
    param_grid_linear,
    cv=5,
    scoring=f05_scorer,
    return_train_score=False,
)
gs_linear_f05.fit(X_train, y_train)

best_linear_f05 = gs_linear_f05.best_estimator_
y_pred_linear_f05 = best_linear_f05.predict(X_test)

print("=== Linear SVC optimized for F0.5 ===")
print("Best params:", gs_linear_f05.best_params_)
print("Test F2 :", round(fbeta_score(y_test, y_pred_linear_f05, beta=2), 4))
print("Test F0.5 :", round(fbeta_score(y_test, y_pred_linear_f05, beta=0.5), 4))

print(
    "\n",
    classification_report(
        y_test, y_pred_linear_f05, target_names=["No Disease", "Disease"]),)

```

```
=== Linear SVC optimized for F0.5 ===  
Best params: {'C': 0.01}  
Test F2      : 0.7778  
Test F0.5    : 0.875
```

	precision	recall	f1-score	support
No Disease	0.81	0.94	0.87	32
Disease	0.91	0.75	0.82	28
accuracy			0.85	60
macro avg	0.86	0.84	0.85	60
weighted avg	0.86	0.85	0.85	60

RBF Kernel

- **gamma** controls how far the influence of a single training point reaches.

High gamma gives a tight localized boundary with overfitting risk, low gamma gives a smoother broader boundary.

```
In [20]: # Hyperparameter grid for RBF kernel  
param_grid_rbf = {"C": [0.01, 0.1, 1, 10, 100], "gamma": [0.001, 0.01, 0.1, 1]}  
  
gs_rbf_f2 = GridSearchCV(  
    estimator=SVC(kernel="rbf", random_state=42),  
    param_grid=param_grid_rbf,  
    scoring=f2_scorer,  
    cv=5,  
    n_jobs=-1,  
    refit=True  
)  
  
# optimize for F2  
gs_rbf_f2.fit(X_train, y_train)  
  
best_rbf_f2 = gs_rbf_f2.best_estimator_  
y_pred_rbf_f2 = best_rbf_f2.predict(X_test)  
  
print("=== RBF optimized for F2 ===")  
print("Best params:", gs_rbf_f2.best_params_)  
print("Test F2      :", round(fbeta_score(y_test, y_pred_rbf_f2, beta=2), 4))  
print("Test F0.5    :", round(fbeta_score(y_test, y_pred_rbf_f2, beta=0.5), 4))  
  
print(  
    "\n",
```

```
classification_report(
    y_test, y_pred_rbf_f2, target_names=["No Disease", "Disease"]),)
```

=== RBF optimized for F2 ===

Best params: {'C': 10, 'gamma': 0.1}

Test F2 : 0.8273

Test F0.5 : 0.8456

	precision	recall	f1-score	support
No Disease	0.85	0.88	0.86	32
Disease	0.85	0.82	0.84	28
accuracy			0.85	60
macro avg	0.85	0.85	0.85	60
weighted avg	0.85	0.85	0.85	60

```
In [21]: gs_rbf_f05 = GridSearchCV(
    estimator=SVC(kernel="rbf", random_state=42),
    param_grid=param_grid_rbf,
    scoring=f05_scorer,
    cv=5,
    n_jobs=-1,
    refit=True
)

# optimize for F0.5
gs_rbf_f05.fit(X_train, y_train)

best_rbf_f05 = gs_rbf_f05.best_estimator_
y_pred_rbf_f05 = best_rbf_f05.predict(X_test)

print("=== RBF optimized for F0.5 ===")
print("Best params:", gs_rbf_f05.best_params_)
print("Test F2 :", round(fbeta_score(y_test, y_pred_rbf_f05, beta=2), 4))
print("Test F0.5 :", round(fbeta_score(y_test, y_pred_rbf_f05, beta=0.5), 4))

print(
    "\n",
    classification_report(
        y_test, y_pred_rbf_f05, target_names=["No Disease", "Disease"]),)
```



```

=== RBF optimized for F0.5 ===
Best params: {'C': 1, 'gamma': 0.01}
Test F2      : 0.8088
Test F0.5    : 0.8871

```

	precision	recall	f1-score	support
No Disease	0.83	0.94	0.88	32
Disease	0.92	0.79	0.85	28
accuracy			0.87	60
macro avg	0.88	0.86	0.86	60
weighted avg	0.87	0.87	0.87	60

Polynomial Kernel

- **degree** controls the complexity of the polynomial decision boundary.

Higher degree means a more complex boundary with overfitting risk

```

In [22]: # Hyperparameter grid for polynomial kernel
param_grid_poly = {"C": [0.01, 0.1, 1, 10, 100], "degree": [2, 3, 4, 5]}

# optimize for F2
gs_poly_f2 = GridSearchCV(
    estimator=SVC(kernel="poly", random_state=42),
    param_grid=param_grid_poly,
    scoring=f2_scorer,
    cv=5,
    n_jobs=-1,
    refit=True
)

gs_poly_f2.fit(X_train, y_train)

best_poly_f2 = gs_poly_f2.best_estimator_
y_pred_poly_f2 = best_poly_f2.predict(X_test)

print("=== Polynomial SVC optimized for F2 ===")
print("Best params:", gs_poly_f2.best_params_)
print("Test F2      :", round(fbeta_score(y_test, y_pred_poly_f2, beta=2), 4))
print("Test F0.5    :", round(fbeta_score(y_test, y_pred_poly_f2, beta=0.5), 4))

print(
    "\n",

```

```
classification_report(
    y_test, y_pred_poly_f2, target_names=["No Disease", "Disease"]),)
```

```
=== Polynomial SVC optimized for F2 ===
Best params: {'C': 10, 'degree': 3}
Test F2      : 0.7609
Test F0.5    : 0.7955
```

	precision	recall	f1-score	support
No Disease	0.79	0.84	0.82	32
Disease	0.81	0.75	0.78	28
accuracy			0.80	60
macro avg	0.80	0.80	0.80	60
weighted avg	0.80	0.80	0.80	60

```
In [23]: # optimize for F0.5
gs_poly_f05 = GridSearchCV(
    estimator=SVC(kernel="poly", random_state=42),
    param_grid=param_grid_poly,
    scoring=f05_scorer,
    cv=5,
    n_jobs=-1,
    refit=True
)

gs_poly_f05.fit(X_train, y_train)

best_poly_f05 = gs_poly_f05.best_estimator_
y_pred_poly_f05 = best_poly_f05.predict(X_test)

print("=== Polynomial SVC optimized for F0.5 ===")
print("Best params:", gs_poly_f05.best_params_)
print("Test F2      :", round(fbeta_score(y_test, y_pred_poly_f05, beta=2), 4))
print("Test F0.5    :", round(fbeta_score(y_test, y_pred_poly_f05, beta=0.5), 4))

print(
    "\n",
    classification_report(
        y_test, y_pred_poly_f05, target_names=["No Disease", "Disease"]),)
```

```

=== Polynomial SVC optimized for F0.5 ===
Best params: {'C': 100, 'degree': 5}
Test F2      : 0.7778
Test F0.5    : 0.875

```

	precision	recall	f1-score	support
No Disease	0.81	0.94	0.87	32
Disease	0.91	0.75	0.82	28
accuracy			0.85	60
macro avg	0.86	0.84	0.85	60
weighted avg	0.86	0.85	0.85	60

Plotting train vs test F2 and F0.5 for all kernels across hyperparameters

```


In [24]: fig, axes = plt.subplots(2, 3, figsize=(18, 10))
fig.suptitle("Train vs Test Score - All Kernels", fontsize=15, fontweight="bold")

kernel_setups = [
    ("Linear", "C", [0.01, 0.1, 1, 10, 100],
     [SVC(kernel="linear", C=c, random_state=42) for c in [0.01, 0.1, 1, 10, 100]]),
    ("RBF", "gamma", [0.001, 0.01, 0.1, 1],
     [SVC(kernel="rbf", C=10, gamma=g, random_state=42) for g in [0.001, 0.01, 0.1, 1]]),
    ("Poly", "degree", [2, 3, 4, 5],
     [SVC(kernel="poly", C=10, degree=d, random_state=42) for d in [2, 3, 4, 5]]),]

for col, (name, param_name, param_vals, svcs) in enumerate(kernel_setups):
    for row, (beta, metric_name) in enumerate([(2, "F2"), (0.5, "F0.5")]):
        train_scores, test_scores = [], []
        for svc in svcs:
            svc.fit(X_train, y_train)
            train_scores.append(round(fbeta_score(y_train, svc.predict(X_train), beta=beta), 4))
            test_scores.append(round(fbeta_score(y_test, svc.predict(X_test), beta=beta), 4))
        ax = axes[row, col]
        ax.plot(param_vals, train_scores, "o-", color="#2196F3", linewidth=2, label="Train")
        ax.plot(param_vals, test_scores, "o-", color="#E91E63", linewidth=2, label="Test")
        ax.fill_between(param_vals, train_scores, test_scores, alpha=0.1, color="gray")
        ax.set_title(f"{name} - {metric_name}", fontsize=12, fontweight="bold")
        ax.set_xlabel(param_name, fontsize=10)
        ax.set_ylabel(f"{metric_name} Score", fontsize=10)
        ax.set_xscale("log" if param_name in ("C", "gamma") else "linear")
        ax.legend(fontsize=10)
        ax.grid(True, linestyle="--", alpha=0.4)
fig.tight_layout()
plt.show()

```

Train vs Test Score - All Kernels



Train vs Test Gap Analysis

Linear Kernel

- F2: Train and test scores are close across all C values (~0.80), confirming **no overfitting**. The model generalizes well, consistent with the dataset being near-linearly separable.
- F0.5: Test exceeds train at low C (C=0.01), indicating **underfitting at high regularization**. Both stabilize around 0.86 after C=0.1, sweet spot is low C.

RBF Kernel

- F2: Train and test track closely at low gamma, but at gamma=1.0 train approaches 1.0 while test collapses to ~0.20 indicating **severe overfitting**. Optimal gamma=0.1 before the divergence.
- F0.5: Same pattern, gamma=1.0 causes the model to memorize training data, destroying generalization. Confirms gamma must stay ≤ 0.1 .

Polynomial Kernel

- F2 & F0.5: Both show a large and growing gap between train and test as degree increases.
- At degree=4, test F2 hits its minimum (~ 0.55) while train stays near 1.0 indicating **maximum overfitting**. degree=3 is the only stable point, which is why GridSearch selected it.
- Polynomial is highly sensitive to degree and prone to overfitting, unlike RBF which degrades only at extreme gamma.

Model Evaluation

```
In [25]: models = [  
    ("Linear", "F2", gs_linear_f2, best_linear_f2, y_pred_linear_f2),  
    ("Linear", "F0.5", gs_linear_f05, best_linear_f05, y_pred_linear_f05),  
    ("RBF", "F2", gs_rbf_f2, best_rbf_f2, y_pred_rbf_f2),  
    ("RBF", "F0.5", gs_rbf_f05, best_rbf_f05, y_pred_rbf_f05),  
    ("Poly", "F2", gs_poly_f2, best_poly_f2, y_pred_poly_f2),  
    ("Poly", "F0.5", gs_poly_f05, best_poly_f05, y_pred_poly_f05)]  
rows = []  
for kernel, metric, gs, _, y_pred in models:  
    rows.append({  
        "Kernel": kernel,  
        "Optimized For": metric,  
        "Best Params": str(gs.best_params_),  
        "Test F2": round(fbeta_score(y_test, y_pred, beta=2), 4),  
        "Test F0.5": round(fbeta_score(y_test, y_pred, beta=0.5), 4)})  
df_results = pd.DataFrame(rows)  
df_results
```

```
Out[25]:
```

	Kernel	Optimized For	Best Params	Test F2	Test F0.5
0	Linear	F2	{'C': 10}	0.8029	0.8594
1	Linear	F0.5	{'C': 0.01}	0.7778	0.8750
2	RBF	F2	{'C': 10, 'gamma': 0.1}	0.8273	0.8456
3	RBF	F0.5	{'C': 1, 'gamma': 0.01}	0.8088	0.8871
4	Poly	F2	{'C': 10, 'degree': 3}	0.7609	0.7955
5	Poly	F0.5	{'C': 100, 'degree': 5}	0.7778	0.8750


Visualizing decision boundaries in PCA-reduced 2D space

```
In [26]: pca_2d = PCA(n_components=2)  
X_train_2d = pca_2d.fit_transform(X_train)
```

```

X_test_2d = pca_2d.transform(X_test)
kernel_configs = [
    ("Linear", SVC(kernel="linear", C=10, random_state=42)),
    ("RBF", SVC(kernel="rbf", C=10, gamma=0.1, random_state=42)),
    ("Poly", SVC(kernel="poly", C=10, degree=3, random_state=42)),]
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
fig.suptitle("SVC Decision Boundaries (PCA 2D)", fontsize=14, fontweight="bold")
x_min, x_max = X_train_2d[:, 0].min() - 0.5, X_train_2d[:, 0].max() + 0.5
y_min, y_max = X_train_2d[:, 1].min() - 0.5, X_train_2d[:, 1].max() + 0.5
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 300),
                     np.linspace(y_min, y_max, 300))
for ax, (name, svc) in zip(axes, kernel_configs):
    svc.fit(X_train_2d, y_train)
    zz = svc.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
    ax.contourf(xx, yy, zz, alpha=0.3, cmap="coolwarm")
    ax.scatter(X_test_2d[:, 0], X_test_2d[:, 1], c=y_test,
               cmap="coolwarm", edgecolors="k", linewidths=0.5, s=50)
    ax.set_title(f"{name} Kernel", fontsize=12, fontweight="bold")
    ax.set_xlabel("PC1"); ax.set_ylabel("PC2")
    ax.grid(True, linestyle="--", alpha=0.4)
fig.tight_layout()
plt.show()

```



Conclusion

- The evaluation results show that the **RBF kernel consistently outperforms the linear and polynomial kernels** across both metrics.
- The *linear kernel provides competitive results, suggesting that the dataset exhibits a partially linear structure*. However, it is consistently outperformed by the RBF kernel, indicating that nonlinear relationships are still present.
- The *polynomial kernel yields the weakest performance across both metrics*, likely due to its sensitivity to the choice of degree and potential overfitting at higher degrees.

Take-home message

Task 1: Clustering

- Choosing the right clustering algorithm matters more than tuning its parameters as *DBSCAN* completely fails here because the data geometry does not match its assumptions.
- *K-Means*, despite its simplicity, can outperform more complex probabilistic models like *GMM* when clusters are compact, spherical, and balanced.
- *Silhouette score* is a reliable unsupervised guide for algorithm selection when ground truth is unavailable.

Task 2: Kernels

- The most complex model is not necessarily the best one for the data, sometimes simple models outperform complex ones.
- There is a trade-off between F2 and F0.5 scores, where optimizing for *F2 minimises false negatives (missed diagnoses)* while *F0.5 minimises false positives (healthy patients incorrectly flagged)*, and the choice should reflect the clinical cost of each error type.